

RateAnalysis

September 16, 2019

1 Introduction

Validating EBDC online compression throughput using the Supermicro SuperWorkstation 7049GP-TRT, with 2 x Intel Xeon Silver 4216 Processor 16-Core 2.1GHz 32 core CPUs and 128 GB memory.

The data is all 2019 sPHENIX TPC SAMPA data at FTBF total 1+TB. The data are buffered on ASUS Hyper M.2 X16 PCIe 3.0 X4 Expansion Card V2 with four SAMSUNG 970 EVO PLUS M.2 2280 1TB PCIe Gen 3.0 x4 NVMe 1.3 V-NAND configured in 4-strip software RAID0. The RAID is tested to 6GBps write and 11GBps write through its PCIe Gen3 x16 interface, matching a large fraction of the FELIX throughput and suppress the expected average rate in sPHENIX year-5 operation.

The data is readout as parallel jobs via [start-compression.sh](#), and sink via TPC connection to multiple ncat processes either at localhost or remote which can be started with [start-sink.sh](#)

2 Inputs

```
[1]: # DataDir = './data_tmp/'  
DataDir = './data_25x_localhost/'  
studytitle = r"$\bf{EBDC}$" + " compression\nlocalhost loopback"  
  
[2]: # %matplotlib widget  
# %matplotlib ipympl  
%matplotlib inline  
# well the html export like dump formats
```

3 Processing

```
[3]: import os  
import ntpath  
import re  
import pandas as pd  
import numpy as np  
  
def processDataset(dataset: str):
```

```

split = dataset.split('-')
if (len(split) != 3):
    print('skip {}'.format (dataset) );
    return;

zipcmd = split[0];
ziplevel = int(split[1]);
jobs = int(split[2]);
datasetDir = os.path.join(DataDir, dataset)

print('processing {}, {} level{} x{}'.format (datasetDir,
→zipcmd,ziplevel,jobs) );

datasubfolders = [os.path.basename(f.path) for f in os.scandir(datasetDir)]
→if f.is_file() ]
datasubfolders.sort()
rpv_in = re.compile('pv_in_([0-9]*)\.log')

for data in datasubfolders:
#     print ('data = {}'.format(data));
    m = rpv_in.search(data)
    if m is not None:
#         print ('found {} -> {}'.format(data, m.group(1)));
        jobID = m.group(1);
        with open(os.path.join(datasetDir, 'pv_in_{}.log'.format(jobID)))
→as f:

            split = f.readlines()[-1].split();
            assert(len(split)==2)
            inTime = float(split[0])
            inSize = float(split[1])
            with open(os.path.join(datasetDir, 'pv_out_{}.log'.format(jobID)))
→as f:

                split = f.readlines()[-1].split();
                assert(len(split)==2)
                outTime = float(split[0])
                outSize = float(split[1])

#         print ('df.append {} . {} , {} -> {}'.format(data,
→jobID,inSize,outSize));
        dictData = { 'dataset' : dataset ,
                    'zipcmd': zipcmd,
                    'ziplevel': ziplevel,
                    'jobs': jobs,
                    'jobID': jobID,
                    'inTime': inTime,
                    'inSize': inSize,
                    'outTime': outTime ,

```

```

        'outSize': outSize
    }
    global dataframe
    dataframe = dataframe.append(dictData, ignore_index=True)

dataframe = pd.DataFrame(columns=['dataset', 'zipcmd', 'ziplevel', 'jobs',
    → 'jobID', 'inTime', 'inSize', 'outTime', 'outSize'])
subfolders = [f.path for f in os.scandir(DataDir) if f.is_dir() ]
subfolders.sort()

for dataset in subfolders:
    processDataset(os.path.basename(dataset))

```

```

processing ./data_tmp/gzip-1-48, gzip level1 x48
processing ./data_tmp/gzip-2-48, gzip level2 x48
processing ./data_tmp/gzip-3-48, gzip level3 x48
processing ./data_tmp/gzip-5-48, gzip level5 x48
processing ./data_tmp/gzip-7-48, gzip level7 x48
processing ./data_tmp/lz4-1-48, lz4 level1 x48
processing ./data_tmp/lz4-2-48, lz4 level2 x48
processing ./data_tmp/lz4-3-48, lz4 level3 x48
processing ./data_tmp/lz4-5-48, lz4 level5 x48
processing ./data_tmp/lz4-7-48, lz4 level7 x48
processing ./data_tmp/lz4-9-48, lz4 level9 x48
processing ./data_tmp/lzop-1-48, lzop level1 x48
processing ./data_tmp/lzop-2-48, lzop level2 x48
processing ./data_tmp/lzop-3-48, lzop level3 x48
processing ./data_tmp/lzop-5-48, lzop level5 x48
processing ./data_tmp/lzop-7-48, lzop level7 x48
processing ./data_tmp/lzop-9-48, lzop level9 x48

```

4 Plot

```

[4]: dataframeSum = pd.DataFrame(columns=['dataset', 'zipcmd', 'ziplevel',
    → 'jobs', 'totalInTime', 'totalInSize', 'totalOutTime', 'totalOutSize',
    → 'Compression', 'inRateGbps', 'outRateGbps'])

zipcmds = dataframe.zipcmd.unique()

for zipcmd in zipcmds:

    zipRows = dataframe.loc[dataframe['zipcmd'] == zipcmd]

    ziplevels = zipRows.ziplevel.unique()

```

```

for ziplevel in ziplevels:
    ziplevelRows = zipRows.loc[zipRows['ziplevel'] == ziplevel]
    print ('processing ', zipcmd, '.',ziplevel, ' size= ',ziplevelRows.
→size, 'compression ratio = ',ziplevelRows['outSize'].sum()/
→ziplevelRows['inSize'].sum())
    assert(ziplevelRows.size>1000)

    dictData = { 'dataset' : ziplevelRows['dataset'].iloc[0] ,
        'zipcmd': ziplevelRows['zipcmd'].iloc[0] ,
        'ziplevel': ziplevelRows['ziplevel'].iloc[0] ,
        'jobs': ziplevelRows['jobs'].iloc[0] ,
        'totalInTime' : ziplevelRows['inTime'].sum() ,
        'totalInSize': ziplevelRows['inSize'].sum() ,
        'totalOutTime': ziplevelRows['outTime'].sum() ,
        'totalOutSize': ziplevelRows['outSize'].sum() ,

        }

    dictData['Compression'] = dictData['totalOutSize']/_
→dictData['totalInSize']
    dictData['inRateGbps'] = dictData['totalInSize']/_
→dictData['totalInTime'] * dictData['jobs'] *8/1e9
    dictData['outRateGbps'] = dictData['totalOutSize']/_
→dictData['totalOutTime']* dictData['jobs'] *8/1e9

    dataframeSum = dataframeSum.append(dictData, ignore_index=True)

```

```

processing gzip . 1 size= 2133 compression ratio = 0.43932139377897234
processing gzip . 2 size= 2133 compression ratio = 0.43639353842830403
processing gzip . 3 size= 2133 compression ratio = 0.424364599250152
processing gzip . 5 size= 2133 compression ratio = 0.4300136029630121
processing gzip . 7 size= 2133 compression ratio = 0.4264271430634125
processing lz4 . 1 size= 2133 compression ratio = 0.6751259046982664
processing lz4 . 2 size= 2133 compression ratio = 0.6751259046982664
processing lz4 . 3 size= 2133 compression ratio = 0.5778558660340661
processing lz4 . 5 size= 2133 compression ratio = 0.5365085644373812
processing lz4 . 7 size= 2133 compression ratio = 0.5207546008082999
processing lz4 . 9 size= 2133 compression ratio = 0.5189943751016245
processing lzop . 1 size= 2133 compression ratio = 0.6377294848760965
processing lzop . 2 size= 2133 compression ratio = 0.6359905949774498
processing lzop . 3 size= 2133 compression ratio = 0.6359905949774498
processing lzop . 5 size= 2133 compression ratio = 0.6359905949774498
processing lzop . 7 size= 2133 compression ratio = 0.48989600087271923
processing lzop . 9 size= 2133 compression ratio = 0.4866962781452209

```

```

[5]: import matplotlib.pyplot as plt
import numpy as np

```

```

Colors = ['#1f77b4',
          '#ff7f0e',
          '#2ca02c',
          '#d62728',
          '#9467bd',
          '#8c564b',
          '#e377c2',
          '#7f7f7f',
          '#bcbd22',
          '#17becf',
          '#1a55ff']
Markers = ['o', 's', 'D', 'p', 'P']

font = {'size' : 14}
plt.rcParams()
plt.rc('font', **font)

studytitle_sup = studytitle + "\n{:d} proc. 2x Xeon4216\n{:.1f}TB TPC FTBF_
→data".format(
    dataframeSum['jobs'].iloc[0], dataframeSum['totalInSize'].iloc[0]/1e12)

```

4.1 Compression plot

```

[6]: # dataframeSum.plot(x = 'ziplevel', y = "Compression")

fig = plt.figure()
ax = fig.add_axes([0.15, 0.15, 0.85, 0.85])
plt.xlabel('Compression level')
plt.ylabel('Compression Ratio [out/in]')


markiter = iter(Markers);
coleriter = iter(Colors);

for zipcmd in dataframeSum.zipcmd.unique():
    zipRows = dataframeSum.loc[dataframeSum['zipcmd'] == zipcmd]
    ax.plot(zipRows['ziplevel'].to_numpy(), zipRows['Compression'].to_numpy(),
            marker=next(markiter), color=next(coleriter), markersize = 8,
→linewidth = 3,
            label=zipcmd)

plt.legend(loc='best', title = studytitle_sup, frameon=False)

plt.savefig(os.path.join(DataDir, "Compression.png"), dpi=150)
plt.savefig(os.path.join(DataDir, "Compression.pdf"), dpi=150)

```



4.2 Compressed throughput

```
[7]: # dataframeSum.plot(x = 'ziplevel', y = "Compression")


fig = plt.figure()
ax = fig.add_axes([0.15, 0.15, 0.85, 0.85])
plt.xlabel('Compression level')
plt.ylabel('Compressed trthroughput [Gbps]')

markiter = iter(Markers);
coleriter = iter(Colors);

for zipcmd in dataframeSum.zipcmd.unique():
    zipRows = dataframeSum.loc[dataframeSum['zipcmd'] == zipcmd]
    ax.plot(zipRows['ziplevel'].to_numpy(), zipRows['outRateGbps'].to_numpy(),
            marker=next(markiter), color=next(coleriter), markersize = 8,
            linewidth = 3,
            label=zipcmd)

plt.legend(loc='best', title = studytitle_sup, frameon=False)
```

```
plt.savefig(os.path.join(DataDir, "Throughput.png"), dpi=150)
plt.savefig(os.path.join(DataDir, "Throughput.pdf"), dpi=150)
```



4.3 Work point curve

```
[8]: # dataframeSum.plot(x = 'ziplevel', y = "Compression")

fig = plt.figure()
ax = fig.add_axes([0.15, 0.15, 0.85, 0.85])
plt.ylabel('Compression ratio [out/in]')
plt.xlabel('Compressed trthroughput [Gbps]')

markiter = iter(Markers);
coleriter = iter(Colors);

for zipcmd in dataframeSum.zipcmd.unique():
    zipRows = dataframeSum.loc[dataframeSum['zipcmd'] == zipcmd]
    outRateGbps = zipRows['outRateGbps'].to_numpy()
```

```

Compression = zipRows['Compression'].to_numpy()
ziplevel = zipRows['ziplevel'].to_numpy()
c = next(coleriter)
ax.plot(outRateGbps, Compression,
        marker=next(marker), color=c, markersize = 8, linewidth = 3,
        label=zipcmd)

for i in range(0, len(outRateGbps)):
    plt.text(outRateGbps[i]+.1, Compression[i]-.01, str(ziplevel[i]),
             → fontsize=9, color=c)

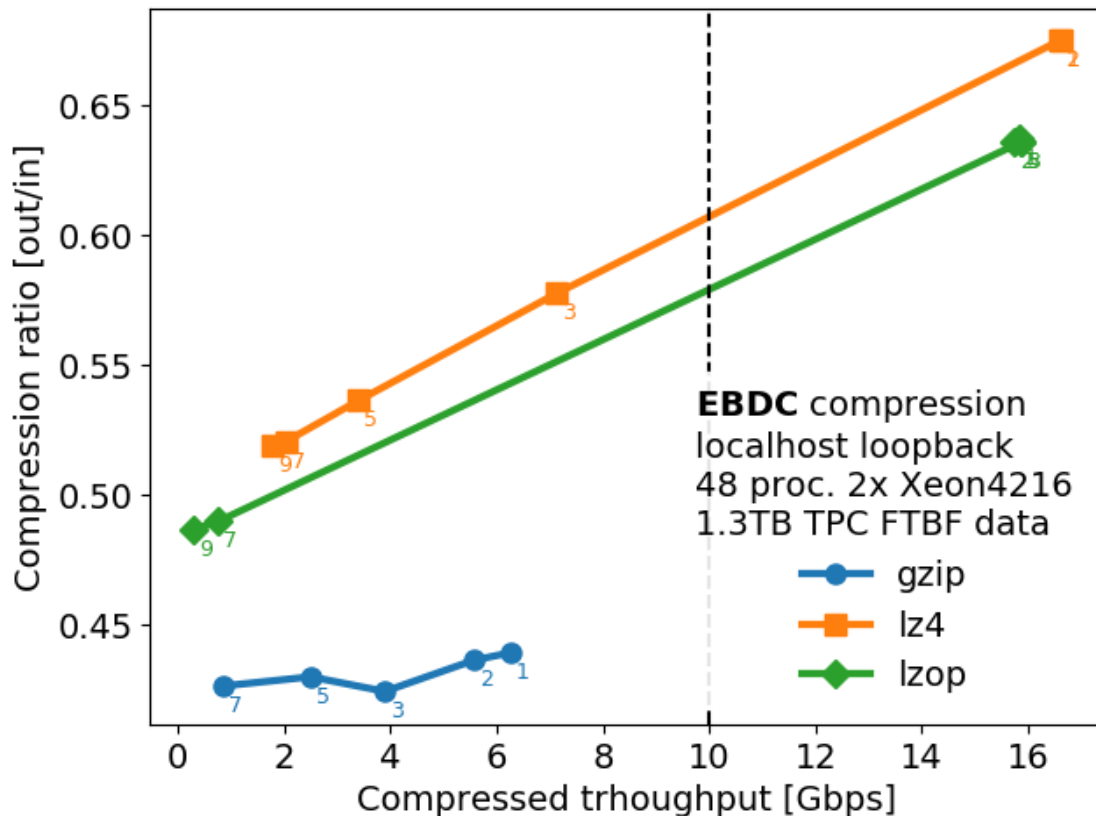
#           [str(i) for i in zipRows['ziplevel'].to_numpy()], fontsize=9)

ax.add_line(plt.Line2D([10, 10], ax.get_ylim(), color = 'black', linestyle =
→ '--'))

plt.legend(loc='best', title = studytitle_sup,
          edgecolor = 'white', frameon=True, facecolor='white', framealpha=0.9)

plt.savefig(os.path.join(DataDir, "FOM.png"), dpi=150)
plt.savefig(os.path.join(DataDir, "FOM.pdf"), dpi=150)

```



5 Scratch

```
[9]: # %save_html os.path.join(DataDir, "analysis.html")
import sys
from subprocess import check_call

d, fname = os.path.split(sys.executable)
# print (d, fname)
check_call([os.path.join(d, 'ipython'), 'nbconvert',
            '--to', 'html',
            'RateAnalysis.ipynb',
            '--output',
            os.path.join(DataDir, "analysis.html")])
check_call([os.path.join(d, 'ipython'), 'nbconvert',
            '--to', 'pdf',
            'RateAnalysis.ipynb',
            '--output',
            os.path.join(DataDir, "analysis.pdf")])
```

[9]: 0

[12]: dataframeSum

[12]:	dataset	zipcmd	ziplevel	jobs	totalInTime	totalInSize	totalOutTime	\
0	gzip-1-48	gzip	1	48	36208.4405	1.345599e+12	36209.0863	
1	gzip-2-48	gzip	2	48	40455.4870	1.345599e+12	40456.2149	
2	gzip-3-48	gzip	3	48	56301.2815	1.345599e+12	56302.1708	
3	gzip-5-48	gzip	5	48	88914.6048	1.345599e+12	88915.8395	
4	gzip-7-48	gzip	7	48	260572.6371	1.345599e+12	260575.8843	
5	lz4-1-48	lz4	1	48	20987.3243	1.345599e+12	20992.1716	
6	lz4-2-48	lz4	2	48	21014.5046	1.345599e+12	21019.3342	
7	lz4-3-48	lz4	3	48	41807.7252	1.345599e+12	41820.8892	
8	lz4-5-48	lz4	5	48	81369.8084	1.345599e+12	81395.9193	
9	lz4-7-48	lz4	7	48	131747.1981	1.345599e+12	131789.8083	
10	lz4-9-48	lz4	9	48	151503.2343	1.345599e+12	151552.0907	
11	lzop-1-48	lzop	1	48	20813.4554	1.345599e+12	20814.1748	
12	lzop-2-48	lzop	2	48	20874.3291	1.345599e+12	20875.2109	
13	lzop-3-48	lzop	3	48	20719.6499	1.345599e+12	20720.4253	
14	lzop-5-48	lzop	5	48	20706.2305	1.345599e+12	20707.0051	
15	lzop-7-48	lzop	7	48	332803.8389	1.345599e+12	332813.0404	
16	lzop-9-48	lzop	9	48	846681.1594	1.345599e+12	846704.0752	
	totalOutSize	Compression	inRateGbps	outRateGbps				
0	5.911506e+11	0.439321	14.270434	6.269195				
1	5.872109e+11	0.436394	12.772313	5.573655				
2	5.710247e+11	0.424365	9.177591	3.894583				

3	5.786260e+11	0.430014	5.811308	2.498907
4	5.738001e+11	0.426427	1.982979	0.845586
5	9.084490e+11	0.675126	24.620107	16.617834
6	9.084490e+11	0.675126	24.588263	16.596359
7	7.775625e+11	0.577856	12.359203	7.139590
8	7.219256e+11	0.536509	6.350146	3.405815
9	7.007271e+11	0.520755	3.921982	2.041730
10	6.983585e+11	0.518994	3.410555	1.769488
11	8.581284e+11	0.637729	24.825775	15.831582
12	8.557885e+11	0.635991	24.753378	15.742251
13	8.557885e+11	0.635991	24.938170	15.859848
14	8.557885e+11	0.635991	24.954333	15.870127
15	6.592038e+11	0.489896	1.552597	0.760590
16	6.548982e+11	0.486696	0.610277	0.297012

[]: