

RateAnalysis

September 16, 2019

1 Introduction

Validating EBDC online compression throughput using the Supermicro SuperWorkstation 7049GP-TRT, with 2 x Intel Xeon Silver 4216 Processor 16-Core 2.1GHz 32 core CPUs and 128 GB memory.

The data is all 2019 sPHENIX TPC SAMPA data at FTBF total 1+TB. The data are buffered on ASUS Hyper M.2 X16 PCIe 3.0 X4 Expansion Card V2 with four SAMSUNG 970 EVO PLUS M.2 2280 1TB PCIe Gen 3.0 x4 NVMe 1.3 V-NAND configured in 4-strip software RAID0. The RAID is tested to 6GBps write and 11GBps write through its PCIe Gen3 x16 interface, matching a large fraction of the FELIX throughput and suppress the expected average rate in sPHENIX year-5 operation.

The data is readout as parallel jobs via [start-compression.sh](#), and sink via TPC connection to multiple ncat processes either at localhost or remote which can be started with [start-sink.sh](#)

2 Inputs

```
[1]: # DataDir = './data_tmp/'
      # DataDir = './data_25x_localhost/'
      DataDir = './data_48x_localhost/'
      studytitle = r"$\bf{EBDC}$" + " compression\nlocalhost loopback"

[2]: # %matplotlib widget
      # %matplotlib ipynb
      %matplotlib inline
      # well the html export like dump formats
```

3 Processing

```
[3]: import os
      import ntpath
      import re
      import pandas as pd
      import numpy as np
```

```

def processDataset(dataset: str):
    split = dataset.split('-')
    if (len(split) != 3):
        print('skip {}'.format (dataset) );
        return;

    zipcmd = split[0];
    ziplevel = int(split[1]);
    jobs = int(split[2]);
    datasetDir = os.path.join(DataDir, dataset)

    print('processing {}, {} level{} x{}'.format (datasetDir,
→zipcmd,ziplevel,jobs) );

    datasubfolders = [os.path.basename(f.path) for f in os.scandir(datasetDir)
→if f.is_file() ]
    datasubfolders.sort()
    rpv_in = re.compile('pv_in_([0-9]*)\.log')

    for data in datasubfolders:
#         print ('data = {}'.format(data));
        m = rpv_in.search(data)
        if m is not None:
#             print ('found {} -> {}'.format(data, m.group(1)));
            jobID = m.group(1);
            with open(os.path.join(datasetDir, 'pv_in_{}.log'.format(jobID)))
→as f:

                split = f.readlines()[-1].split();
                assert(len(split)==2)
                inTime = float(split[0])
                inSize = float(split[1])
                with open(os.path.join(datasetDir, 'pv_out_{}.log'.format(jobID)))
→as f:

                    split = f.readlines()[-1].split();
                    assert(len(split)==2)
                    outTime = float(split[0])
                    outSize = float(split[1])

#             print ('df.append {} . {} , {} -> {}'.format(data,
→jobID,inSize,outSize));
            dictData = { 'dataset' : dataset ,
                        'zipcmd': zipcmd,
                        'ziplevel': ziplevel,
                        'jobs': jobs,
                        'jobID': jobID,
                        'inTime': inTime,
                        'inSize': inSize,

```

```

        'outTime': outTime ,
        'outSize': outSize
    }
    global dataframe
    dataframe = dataframe.append(dictData, ignore_index=True)

dataframe = pd.DataFrame(columns=['dataset', 'zipcmd', 'ziplevel', 'jobs',
    ↪ 'jobID', 'inTime', 'inSize', 'outTime', 'outSize'])
subfolders = [f.path for f in os.scandir(DataDir) if f.is_dir() ]
subfolders.sort()

for dataset in subfolders:
    processDataset(os.path.basename(dataset))

```

```

skip .ipynb_checkpoints
processing ./data_25x_localhost/gzip-1-25, gzip level1 x25
processing ./data_25x_localhost/gzip-2-25, gzip level2 x25
processing ./data_25x_localhost/gzip-3-25, gzip level3 x25
processing ./data_25x_localhost/gzip-5-25, gzip level5 x25
processing ./data_25x_localhost/gzip-7-25, gzip level7 x25
processing ./data_25x_localhost/gzip-9-25, gzip level9 x25
processing ./data_25x_localhost/lz4-1-25, lz4 level1 x25
processing ./data_25x_localhost/lz4-2-25, lz4 level2 x25
processing ./data_25x_localhost/lz4-3-25, lz4 level3 x25
processing ./data_25x_localhost/lz4-5-25, lz4 level5 x25
processing ./data_25x_localhost/lz4-7-25, lz4 level7 x25
processing ./data_25x_localhost/lz4-9-25, lz4 level9 x25
processing ./data_25x_localhost/lzop-1-25, lzop level1 x25
processing ./data_25x_localhost/lzop-2-25, lzop level2 x25
processing ./data_25x_localhost/lzop-3-25, lzop level3 x25
processing ./data_25x_localhost/lzop-5-25, lzop level5 x25
processing ./data_25x_localhost/lzop-7-25, lzop level7 x25
processing ./data_25x_localhost/lzop-9-25, lzop level9 x25

```

4 Plot

```

[4]: dataframeSum = pd.DataFrame(columns=['dataset', 'zipcmd', 'ziplevel',
    ↪ 'jobs', 'totalInTime', 'totalInSize', 'totalOutTime', 'totalOutSize',
    ↪ 'Compression', 'inRateGbps', 'outRateGbps'])

zipcmds = dataframe.zipcmd.unique()

for zipcmd in zipcmds:

```

```

zipRows = dataframe.loc[dataframe['zipcmd'] == zipcmd]

ziplevels = zipRows.ziplevel.unique()
for ziplevel in ziplevels:
    ziplevelRows = zipRows.loc[zipRows['ziplevel'] == ziplevel]
    print ('processing ', zipcmd, '.',ziplevel, ' size= ',ziplevelRows.
→size, 'compression ratio = ',ziplevelRows['outSize'].sum()/
→ziplevelRows['inSize'].sum())
    assert(ziplevelRows.size>1000)

    dictData = { 'dataset' : ziplevelRows['dataset'].iloc[0] ,
                  'zipcmd': ziplevelRows['zipcmd'].iloc[0] ,
                  'ziplevel': ziplevelRows['ziplevel'].iloc[0] ,
                  'jobs': ziplevelRows['jobs'].iloc[0] ,
                  'totalInTime' : ziplevelRows['inTime'].sum() ,
                  'totalInSize': ziplevelRows['inSize'].sum() ,
                  'totalOutTime': ziplevelRows['outTime'].sum() ,
                  'totalOutSize': ziplevelRows['outSize'].sum() ,

                }

    dictData['Compression'] = dictData['totalOutSize']/_
→dictData['totalInSize']
    dictData['inRateGbps'] = dictData['totalInSize']/_
→dictData['totalInTime'] * dictData['jobs'] *8/1e9
    dictData['outRateGbps'] = dictData['totalOutSize']/_
→dictData['totalOutTime']* dictData['jobs'] *8/1e9

    dataframeSum = dataframeSum.append(dictData, ignore_index=True)

```

```

processing gzip . 1 size= 2133 compression ratio = 0.43932139377897234
processing gzip . 2 size= 2133 compression ratio = 0.43639353842830403
processing gzip . 3 size= 2133 compression ratio = 0.424364599250152
processing gzip . 5 size= 2133 compression ratio = 0.4300136029630121
processing gzip . 7 size= 2133 compression ratio = 0.4264271430634125
processing gzip . 9 size= 2133 compression ratio = 0.42370792929411943
processing lz4 . 1 size= 2133 compression ratio = 0.6751259046982664
processing lz4 . 2 size= 2133 compression ratio = 0.6751259046982664
processing lz4 . 3 size= 2133 compression ratio = 0.5778558660340661
processing lz4 . 5 size= 2133 compression ratio = 0.5365085644373812
processing lz4 . 7 size= 2133 compression ratio = 0.5207546008082999
processing lz4 . 9 size= 2133 compression ratio = 0.5189943751016245
processing lzop . 1 size= 2133 compression ratio = 0.6377294848760965
processing lzop . 2 size= 2133 compression ratio = 0.6359905949774498
processing lzop . 3 size= 2133 compression ratio = 0.6359905949774498
processing lzop . 5 size= 2133 compression ratio = 0.6359905949774498
processing lzop . 7 size= 2133 compression ratio = 0.48989600087271923

```

```
processing lzop . 9 size= 2133 compression ratio = 0.4866962781452209
```

```
[5]: import matplotlib.pyplot as plt
import numpy as np

Colors = ['#1f77b4',
          '#ff7f0e',
          '#2ca02c',
          '#d62728',
          '#9467bd',
          '#8c564b',
          '#e377c2',
          '#7f7f7f',
          '#bcbd22',
          '#17becf',
          '#1a55ff']

Markers = ['o' , 's' , 'D' , 'p' , 'P']

font = {'size' : 14}
plt.rcParams()
plt.rc('font', **font)

studytitle_sup = studytitle + "\n{:d} proc. 2x Xeon4216\n{:.1f}TB TPC FTBF_
→data".format(
    dataframeSum['jobs'].iloc[0], dataframeSum['totalInSize'].iloc[0]/1e12)
```

4.1 Compression plot

```
[11]: # dataframeSum.plot(x = 'ziplevel', y = "Compression")


fig = plt.figure()
ax = fig.add_axes([0.15, 0.15, 0.85, 0.85])
plt.xlabel('Compression level')
plt.ylabel('Compression Ratio [out/in]')

markiter = iter(Markers);
coleriter = iter(Colors);

for zipcmd in dataframeSum.zipcmd.unique():
    zipRows = dataframeSum.loc[dataframeSum['zipcmd'] == zipcmd]
    ax.plot(zipRows['ziplevel'].to_numpy(), zipRows['Compression'].to_numpy(),
            marker=next(markiter), color=next(coleriter), markersize = 8,
→linewidth = 3,
            label=zipcmd)

plt.legend(loc='best', title = studytitle_sup, frameon=False)
```

```
plt.savefig(os.path.join(DataDir, "Compression.png"), dpi=150)
plt.savefig(os.path.join(DataDir, "Compression.pdf"), dpi=150)
```



4.2 Compressed throughput

```
[7]: # dataframeSum.plot(x = 'ziplevel', y = "Compression")

fig = plt.figure()
ax = fig.add_axes([0.15, 0.15, 0.85, 0.85])
plt.xlabel('Compression level')
plt.ylabel('Compressed throughput [Gbps]')

markiter = iter(Markers);
coleriter = iter(Colors);

for zipcmd in dataframeSum.zipcmd.unique():
    zipRows = dataframeSum.loc[dataframeSum['zipcmd'] == zipcmd]
    ax.plot(zipRows['ziplevel'].to_numpy(), zipRows['outRateGbps'].to_numpy(),
            marker=next(markiter), color=next(coleriter), markersize = 8,
            linewidth = 3,
```


```

label=zipcmd)

plt.legend(loc='best',title = studytitle_sup, frameon=False)

plt.savefig(os.path.join(DataDir,"Throughput.png"), dpi=150)
plt.savefig(os.path.join(DataDir,"Throughput.pdf"), dpi=150)

```



4.3 Work point curve

```

[8]: # dataframeSum.plot(x = 'ziplevel', y = "Compression")

fig = plt.figure()
ax = fig.add_axes([0.15, 0.15, 0.85, 0.85])
plt.ylabel('Compression ratio [out/in]')
plt.xlabel('Compressed trthroughput [Gbps]')

markiter = iter(Markers);
coleriter = iter(Colors);

```

```

for zipcmd in dataframeSum.zipcmd.unique():
    zipRows = dataframeSum.loc[dataframeSum['zipcmd'] == zipcmd]
    outRateGbps = zipRows['outRateGbps'].to_numpy()
    Compression = zipRows['Compression'].to_numpy()
    ziplevel = zipRows['ziplevel'].to_numpy()
    c = next(coleriter)
    ax.plot(outRateGbps, Compression,
            marker=next(markiter), color=c, markersize = 8, linewidth = 3,
            label=zipcmd)

    for i in range(0, len(outRateGbps)):
        plt.text(outRateGbps[i]+.1, Compression[i]-.01, str(ziplevel[i]),
        ↪ fontsize=9, color=c)


#             [str(i) for i in zipRows['ziplevel'].to_numpy()], fontsize=9)

ax.add_line(plt.Line2D([10, 10], ax.get_ylim(), color = 'black', linestyle =
    ↪ '--'))

plt.legend(loc='best', title = studytitle_sup,
          edgecolor = 'white', frameon=True, facecolor='white', framealpha=0.9)

plt.savefig(os.path.join(DataDir, "FOM.png"), dpi=150)
plt.savefig(os.path.join(DataDir, "FOM.pdf"), dpi=150)

```

5 Scratch

```
[9]: # %save_html os.path.join(DataDir, "analysis.html")
import sys
from subprocess import check_call

d, fname = os.path.split(sys.executable)
# print (d, fname)
check_call([os.path.join(d, 'ipython'), 'nbconvert',
              '--to', 'html',
              'RateAnalysis.ipynb',
              '--output',
              os.path.join(DataDir, "analysis.html")])
check_call([os.path.join(d, 'ipython'), 'nbconvert',
              '--to', 'pdf',
              'RateAnalysis.ipynb',
              '--output',
              os.path.join(DataDir, "analysis.pdf")])
```

[9]: 0

```
[10]: dataframeSum
```

```
[10]:      dataset zipcmd ziplevel jobs totalInTime totalInSize totalOutTime \
0  gzip-1-25  gzip      1    25   27710.4416  1.345599e+12   27711.0109
1  gzip-2-25  gzip      2    25   31493.8780  1.345599e+12   31494.5092
2  gzip-3-25  gzip      3    25   46565.4287  1.345599e+12   46566.2482
3  gzip-5-25  gzip      5    25   73615.0219  1.345599e+12   73616.1849
4  gzip-7-25  gzip      7    25  225600.2495  1.345599e+12  225603.1836
5  gzip-9-25  gzip      9    25  432304.2058  1.345599e+12  432309.5521
6  lz4-1-25   lz4       1    25   12061.4768  1.345599e+12   12064.6152
7  lz4-2-25   lz4       2    25   12057.8337  1.345599e+12   12060.9059
8  lz4-3-25   lz4       3    25   31553.6753  1.345599e+12   31563.7981
9  lz4-5-25   lz4       5    25   62380.4076  1.345599e+12   62400.6044
10 lz4-7-25   lz4       7    25  103179.2939  1.345599e+12  103212.2861
11 lz4-9-25   lz4       9    25  120706.8721  1.345599e+12  120745.8980
12 lzop-1-25  lzop       1    25   11651.7342  1.345599e+12   11652.2921
13 lzop-2-25  lzop       2    25   11740.9330  1.345599e+12   11741.5004
14 lzop-3-25  lzop       3    25   11739.2630  1.345599e+12   11739.8147
15 lzop-5-25  lzop       5    25   11729.0190  1.345599e+12   11729.5595
16 lzop-7-25  lzop       7    25  274774.6620  1.345599e+12  274782.5120
17 lzop-9-25  lzop       9    25  754100.0086  1.345599e+12  754120.6388
```

	totalOutSize	Compression	inRateGbps	outRateGbps
0	5.911506e+11	0.439321	9.711858	4.266539
1	5.872109e+11	0.436394	8.545149	3.728973
2	5.710247e+11	0.424365	5.779392	2.452526
3	5.786260e+11	0.430014	3.655774	1.572008
4	5.738001e+11	0.426427	1.192906	0.508681
5	5.701411e+11	0.423708	0.622524	0.263765
6	9.084490e+11	0.675126	22.312349	15.059726
7	9.084490e+11	0.675126	22.319090	15.064358
8	7.775625e+11	0.577856	8.528955	4.926926
9	7.219256e+11	0.536509	4.314173	2.313842
10	7.007271e+11	0.520755	2.608274	1.357837
11	6.983585e+11	0.518994	2.229532	1.156741
12	8.581284e+11	0.637729	23.096980	14.728920
13	8.557885e+11	0.635991	22.921507	14.577158
14	8.557885e+11	0.635991	22.924768	14.579251
15	8.557885e+11	0.635991	22.944790	14.591998
16	6.592038e+11	0.489896	0.979420	0.479800
17	6.548982e+11	0.486696	0.356876	0.173685

```
[ ]:
```