

RateAnalysis

September 17, 2019

1 Introduction

Validating EBDC online compression throughput using the Supermicro SuperWorkstation 7049GP-TRT, with 2 x Intel Xeon Silver 4216 Processor 16-Core 2.1GHz 32 core CPUs and 128 GB memory.

The data is all 2019 sPHENIX TPC SAMPA data at FTBF total 1+TB. The data are buffered on ASUS Hyper M.2 X16 PCIe 3.0 X4 Expansion Card V2 with four SAMSUNG 970 EVO PLUS M.2 2280 1TB PCIe Gen 3.0 x4 NVMe 1.3 V-NAND configured in 4-strip software RAID0. The RAID is tested to 6GBps write and 11GBps write through its PCIe Gen3 x16 interface, matching a large fraction of the FELIX throughput and suppress the expected average rate in sPHENIX year-5 operation.

The data is readout as parallel jobs via [start-compression.sh](#), and sink via either /dev/null or TPC connections to multiple ncat processes either at localhost or remote which can be started with [start-sink.sh](#)

2 Inputs

```
[1]: # DataDir = './data_tmp/'  
# DataDir = './data_25x_localhost/'  
# DataDir = './data_48x_localhost/'  
# studytitle = r"${bf{EBDC}}$" + " compression\nlocalhost loopback"  
DataDir = './data_48x_null/'  
studytitle = r"${bf{EBDC}}$" + " compression\noutput to /dev/null"
```

```
[2]: # %matplotlib widget  
# %matplotlib ipympl  
%matplotlib inline  
# well the html export like dump formats
```

3 Processing

```
[3]: import os  
import ntpath  
import re  
import pandas as pd
```

```

import numpy as np

def processDataset(dataset: str):
    split = dataset.split('-')
    if (len(split) != 3):
        print('skip {}'.format (dataset) );
        return;

    zipcmd = split[0];
    ziplevel = int(split[1]);
    jobs = int(split[2]);
    datasetDir = os.path.join(DataDir, dataset)

    print('processing {}, {} level{} x{}'.format (datasetDir,
→zipcmd,ziplevel,jobs) );

    datasubfolders = [os.path.basename(f.path) for f in os.scandir(datasetDir)
→if f.is_file() ]
    datasubfolders.sort()
    rpv_in = re.compile('pv_in_([0-9]*)\.log')

    for data in datasubfolders:
#         print ('data = {}'.format(data));
        m = rpv_in.search(data)
        if m is not None:
#             print ('found {} -> {}'.format(data, m.group(1)));
            jobID = m.group(1);
            with open(os.path.join(datasetDir, 'pv_in_{}.log'.format(jobID)))
→as f:

                split = f.readlines()[-1].split();
                assert(len(split)==2)
                inTime = float(split[0])
                inSize = float(split[1])
                with open(os.path.join(datasetDir, 'pv_out_{}.log'.format(jobID)))
→as f:

                    split = f.readlines()[-1].split();
                    assert(len(split)==2)
                    outTime = float(split[0])
                    outSize = float(split[1])

#             print ('df.append {} . {} , {} -> {}'.format(data,
→jobID,inSize,outSize));
            dictData = { 'dataset' : dataset ,
                        'zipcmd': zipcmd,
                        'ziplevel': ziplevel,
                        'jobs': jobs,

```

```

        'jobID': jobID,
        'inTime': inTime,
        'inSize': inSize,
        'outTime': outTime ,
        'outSize': outSize
    }

    global dataframe
    dataframe = dataframe.append(dictData, ignore_index=True)

dataframe = pd.DataFrame(columns=['dataset', 'zipcmd', 'ziplevel', 'jobs',
    → 'jobID', 'inTime', 'inSize', 'outTime', 'outSize'])
subfolders = [f.path for f in os.scandir(DataDir) if f.is_dir() ]
subfolders.sort()

for dataset in subfolders:
    processDataset(os.path.basename(dataset))

```

```

processing ./data_48x_localhost/gzip-1-48, gzip level1 x48
processing ./data_48x_localhost/gzip-2-48, gzip level2 x48
processing ./data_48x_localhost/gzip-3-48, gzip level3 x48
processing ./data_48x_localhost/gzip-5-48, gzip level5 x48
processing ./data_48x_localhost/gzip-7-48, gzip level7 x48
processing ./data_48x_localhost/gzip-9-48, gzip level9 x48
processing ./data_48x_localhost/lz4-1-48, lz4 level1 x48
processing ./data_48x_localhost/lz4-2-48, lz4 level2 x48
processing ./data_48x_localhost/lz4-3-48, lz4 level3 x48
processing ./data_48x_localhost/lz4-5-48, lz4 level5 x48
processing ./data_48x_localhost/lz4-7-48, lz4 level7 x48
processing ./data_48x_localhost/lz4-9-48, lz4 level9 x48
processing ./data_48x_localhost/lzop-1-48, lzop level1 x48
processing ./data_48x_localhost/lzop-2-48, lzop level2 x48
processing ./data_48x_localhost/lzop-3-48, lzop level3 x48
processing ./data_48x_localhost/lzop-5-48, lzop level5 x48
processing ./data_48x_localhost/lzop-7-48, lzop level7 x48
processing ./data_48x_localhost/lzop-9-48, lzop level9 x48

```

4 Plot

```

[4]: dataframeSum = pd.DataFrame(columns=['dataset', 'zipcmd', 'ziplevel',
    → 'jobs', 'totalInTime', 'totalInSize', 'totalOutTime', 'totalOutSize',
    → 'Compression', 'inRateGbps', 'outRateGbps'])

zipcmds = dataframe.zipcmd.unique()

```

```

for zipcmd in zipcmds:

    zipRows = dataframe.loc[dataframe['zipcmd'] == zipcmd]

    ziplevels = zipRows.ziplevel.unique()
    for ziplevel in ziplevels:
        ziplevelRows = zipRows.loc[zipRows['ziplevel'] == ziplevel]
        print ('processing ', zipcmd, '.', ziplevel, ' size= ', ziplevelRows.
→size, 'compression ratio = ', ziplevelRows['outSize'].sum()/
→ziplevelRows['inSize'].sum())
        assert(ziplevelRows.size>1000)

        dictData = { 'dataset' : ziplevelRows['dataset'].iloc[0] ,
                      'zipcmd': ziplevelRows['zipcmd'].iloc[0] ,
                      'ziplevel': ziplevelRows['ziplevel'].iloc[0] ,
                      'jobs': ziplevelRows['jobs'].iloc[0] ,
                      'totalInTime' : ziplevelRows['inTime'].sum() ,
                      'totalInSize': ziplevelRows['inSize'].sum() ,
                      'totalOutTime': ziplevelRows['outTime'].sum() ,
                      'totalOutSize': ziplevelRows['outSize'].sum() ,

                      }

        dictData['Compression'] = dictData['totalOutSize']/_
→dictData['totalInSize']
        dictData['inRateGbps'] = dictData['totalInSize']/_
→dictData['totalInTime'] * dictData['jobs'] *8/1e9
        dictData['outRateGbps'] = dictData['totalOutSize']/_
→dictData['totalOutTime']* dictData['jobs'] *8/1e9

        dataframeSum = dataframeSum.append(dictData, ignore_index=True)

```

```

processing gzip . 1 size= 2133 compression ratio = 0.43932139377897234
processing gzip . 2 size= 2133 compression ratio = 0.43639353842830403
processing gzip . 3 size= 2133 compression ratio = 0.424364599250152
processing gzip . 5 size= 2133 compression ratio = 0.4300136029630121
processing gzip . 7 size= 2133 compression ratio = 0.4264271430634125
processing gzip . 9 size= 2133 compression ratio = 0.42370792929411943
processing lz4 . 1 size= 2133 compression ratio = 0.6751259046982664
processing lz4 . 2 size= 2133 compression ratio = 0.6751259046982664
processing lz4 . 3 size= 2133 compression ratio = 0.5778558660340661
processing lz4 . 5 size= 2133 compression ratio = 0.5365085644373812
processing lz4 . 7 size= 2133 compression ratio = 0.5207546008082999
processing lz4 . 9 size= 2133 compression ratio = 0.5189943751016245
processing lzop . 1 size= 2133 compression ratio = 0.6377294848760965
processing lzop . 2 size= 2133 compression ratio = 0.6359905949774498
processing lzop . 3 size= 2133 compression ratio = 0.6359905949774498

```

```
processing lzop . 5 size= 2133 compression ratio = 0.6359905949774498
processing lzop . 7 size= 2133 compression ratio = 0.48989600087271923
processing lzop . 9 size= 2133 compression ratio = 0.4866962781452209
```

```
[5]: import matplotlib.pyplot as plt
import numpy as np

Colors = ['#1f77b4',
          '#ff7f0e',
          '#2ca02c',
          '#d62728',
          '#9467bd',
          '#8c564b',
          '#e377c2',
          '#7f7f7f',
          '#bcbd22',
          '#17becf',
          '#1a55ff']
Markers = ['o', 's', 'D', 'p', 'P']

font = {'size' : 14}
plt.rcdefaults()
plt.rc('font', **font)

studytitle_sup = studytitle + "\n{:d} proc. 2x Xeon4216\n{:.1f}TB TPC FTBF_
→data".format(
    dataframeSum['jobs'].iloc[0], dataframeSum['totalInSize'].iloc[0]/1e12)
```

4.1 Compression plot

```
[6]: # dataframeSum.plot(x = 'ziplevel', y = "Compression")

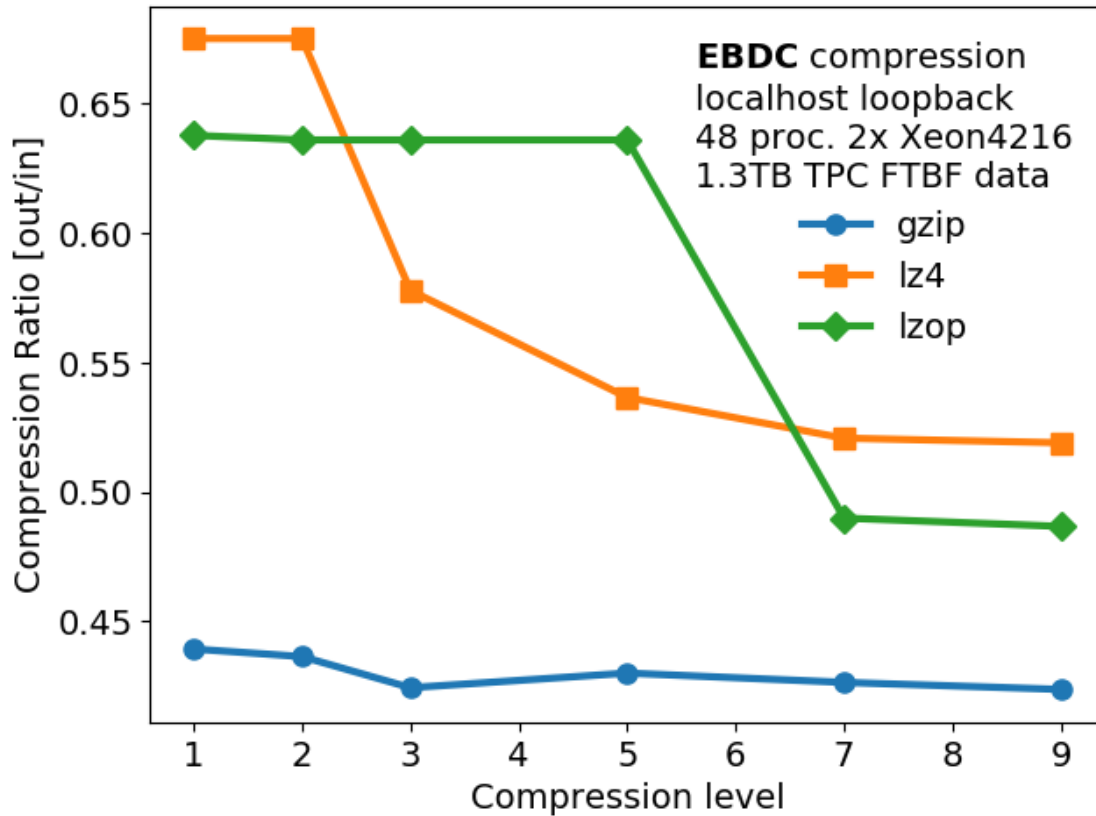
fig = plt.figure()
ax = fig.add_axes([0.15, 0.15, 0.85, 0.85])
plt.xlabel('Compression level')
plt.ylabel('Compression Ratio [out/in]')

markiter = iter(Markers);
coleriter = iter(Colors);

for zipcmd in dataframeSum.zipcmd.unique():
    zipRows = dataframeSum.loc[dataframeSum['zipcmd'] == zipcmd]
    ax.plot(zipRows['ziplevel'].to_numpy(), zipRows['Compression'].to_numpy(),
            marker=next(markiter), color=next(coleriter), markersize = 8,
            →linewidth = 3,
            label=zipcmd)
```

```
plt.legend(loc='best',title = studytitle_sup, frameon=False)

plt.savefig(os.path.join(DataDir,"Compression.png"), dpi=150)
plt.savefig(os.path.join(DataDir,"Compression.pdf"), dpi=150)
```



4.2 Compressed throughput

```
[7]: # dataframeSum.plot(x = 'ziplevel', y = "Compression")

fig = plt.figure()
ax = fig.add_axes([0.15, 0.15, 0.85, 0.85])
plt.xlabel('Compression level')
plt.ylabel('Compressed throughput [Gbps]')

markiter = iter(Markers);
coleriter = iter(Colors);

for zipcmd in dataframeSum.zipcmd.unique():
    zipRows = dataframeSum.loc[dataframeSum['zipcmd'] == zipcmd]
    ax.plot(zipRows['ziplevel'].to_numpy(), zipRows['outRateGbps'].to_numpy(),
```


```

        marker=next(markeriter), color=next(coleriter), markersize = 8,
        linewidth = 3,
        label=zipcmd)

plt.legend(loc='best',title = studytitle_sup, frameon=False)

plt.savefig(os.path.join(DataDir,"Throughput.png"), dpi=150)
plt.savefig(os.path.join(DataDir,"Throughput.pdf"), dpi=150)

```



4.3 Work point curve

```

[8]: # dataframeSum.plot(x = 'ziplevel', y = "Compression")

fig = plt.figure()
ax = fig.add_axes([0.15, 0.15, 0.85, 0.85])
plt.ylabel('Compression ratio [out/in]')
plt.xlabel('Compressed throughput [Gbps]')

markiter = iter(Markers);

```

```

coleriter = iter(Colors);

for zipcmd in dataframeSum.zipcmd.unique():
    zipRows = dataframeSum.loc[dataframeSum['zipcmd'] == zipcmd]
    outRateGbps = zipRows['outRateGbps'].to_numpy()
    Compression = zipRows['Compression'].to_numpy()
    ziplevel = zipRows['ziplevel'].to_numpy()
    c = next(coleriter)
    ax.plot(outRateGbps, Compression,
            marker=next(markiter), color=c, markersize = 8, linewidth = 3,
            label=zipcmd)

    for i in range(0, len(outRateGbps)):
        plt.text(outRateGbps[i]+.1, Compression[i]-.01, str(ziplevel[i]),
        ↪ fontsize=9, color=c)

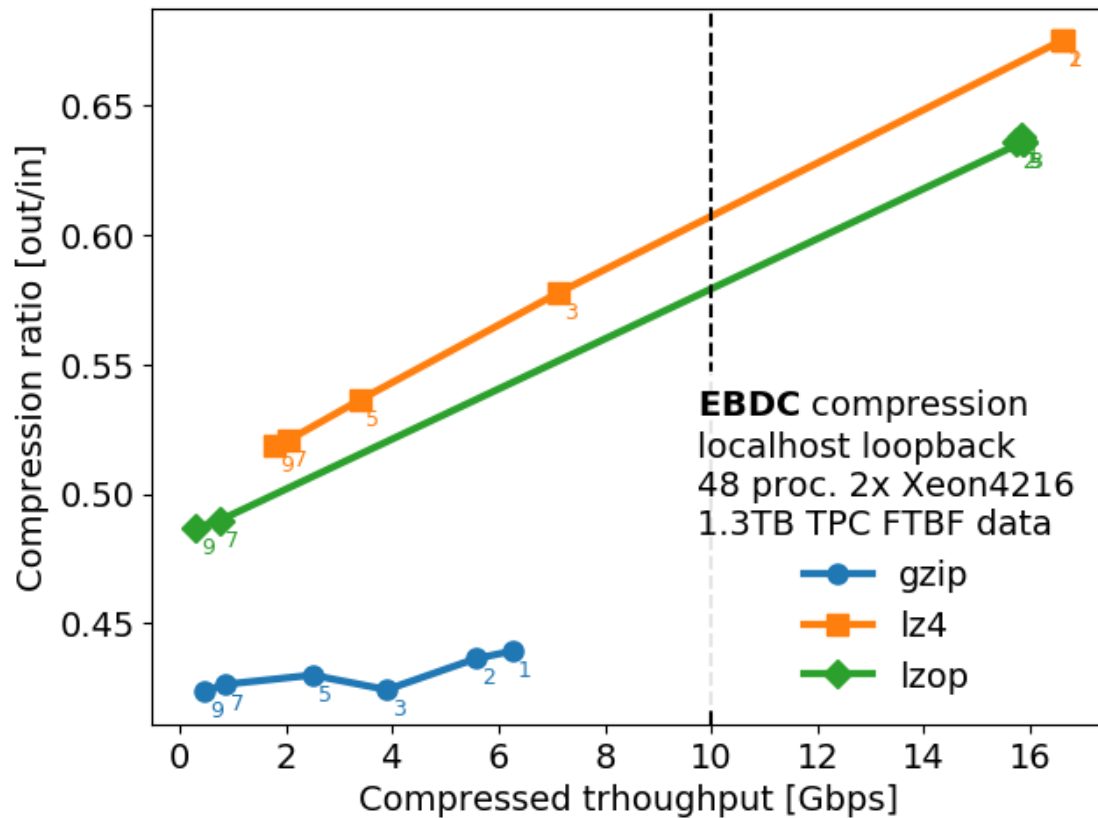
#             [str(i) for i in zipRows['ziplevel'].to_numpy()], fontsize=9)

ax.add_line(plt.Line2D([10, 10], ax.get_ylim(), color = 'black', linestyle =
↪ '--'))

plt.legend(loc='best', title = studytitle_sup,
          edgecolor = 'white', frameon=True, facecolor='white', framealpha=0.9)

plt.savefig(os.path.join(DataDir, "FOM.png"), dpi=150)
plt.savefig(os.path.join(DataDir, "FOM.pdf"), dpi=150)

```

5 Scratch

```
[9]: # %save_html os.path.join(DataDir, "analysis.html")
import sys
from subprocess import check_call

d, fname = os.path.split(sys.executable)
# print (d, fname)
check_call([os.path.join(d, 'ipython'), 'nbconvert',
              '--to', 'html',
              'RateAnalysis.ipynb',
              '--output',
              os.path.join(DataDir, "analysis.html")])
check_call([os.path.join(d, 'ipython'), 'nbconvert',
              '--to', 'pdf',
              'RateAnalysis.ipynb',
              '--output',
              os.path.join(DataDir, "analysis.pdf")])
```

[9]: 0

```
[10]: dataframeSum
```

```
[10]:      dataset zipcmd ziplevel jobs totalInTime totalInSize totalOutTime \
0  gzip-1-48  gzip      1    48   36208.4405  1.345599e+12   36209.0863
1  gzip-2-48  gzip      2    48   40455.4870  1.345599e+12   40456.2149
2  gzip-3-48  gzip      3    48   56301.2815  1.345599e+12   56302.1708
3  gzip-5-48  gzip      5    48   88914.6048  1.345599e+12   88915.8395
4  gzip-7-48  gzip      7    48  260572.6371  1.345599e+12  260575.8843
5  gzip-9-48  gzip      9    48  469586.5532  1.345599e+12  469592.2606
6  lz4-1-48   lz4       1    48   20987.3243  1.345599e+12   20992.1716
7  lz4-2-48   lz4       2    48   21014.5046  1.345599e+12   21019.3342
8  lz4-3-48   lz4       3    48   41807.7252  1.345599e+12   41820.8892
9  lz4-5-48   lz4       5    48   81369.8084  1.345599e+12   81395.9193
10 lz4-7-48   lz4       7    48  131747.1981  1.345599e+12  131789.8083
11 lz4-9-48   lz4       9    48  151503.2343  1.345599e+12  151552.0907
12 lzop-1-48  lzop       1    48   20813.4554  1.345599e+12   20814.1748
13 lzop-2-48  lzop       2    48   20874.3291  1.345599e+12   20875.2109
14 lzop-3-48  lzop       3    48   20719.6499  1.345599e+12   20720.4253
15 lzop-5-48  lzop       5    48   20706.2305  1.345599e+12   20707.0051
16 lzop-7-48  lzop       7    48  332803.8389  1.345599e+12  332813.0404
17 lzop-9-48  lzop       9    48  846681.1594  1.345599e+12  846704.0752
```

	totalOutSize	Compression	inRateGbps	outRateGbps
0	5.911506e+11	0.439321	14.270434	6.269195
1	5.872109e+11	0.436394	12.772313	5.573655
2	5.710247e+11	0.424365	9.177591	3.894583
3	5.786260e+11	0.430014	5.811308	2.498907
4	5.738001e+11	0.426427	1.982979	0.845586
5	5.701411e+11	0.423708	1.100351	0.466222
6	9.084490e+11	0.675126	24.620107	16.617834
7	9.084490e+11	0.675126	24.588263	16.596359
8	7.775625e+11	0.577856	12.359203	7.139590
9	7.219256e+11	0.536509	6.350146	3.405815
10	7.007271e+11	0.520755	3.921982	2.041730
11	6.983585e+11	0.518994	3.410555	1.769488
12	8.581284e+11	0.637729	24.825775	15.831582
13	8.557885e+11	0.635991	24.753378	15.742251
14	8.557885e+11	0.635991	24.938170	15.859848
15	8.557885e+11	0.635991	24.954333	15.870127
16	6.592038e+11	0.489896	1.552597	0.760590
17	6.548982e+11	0.486696	0.610277	0.297012

```
[ ]:
```